

Chapter 6 Modeling users

In [Chapter 5](#), we ended with a stub page for creating new users ([Section 5.4](#)). Over the course of the next six chapters, we'll fulfill the promise implicit in this incipient signup page. In this chapter, we'll take the first critical step by creating a *data model* for users of our site, together with a way to store that data. In [Chapter 7](#), we'll give users the ability to sign up for our site and create a user profile page. Once users can sign up, we'll let them log in and log out as well ([Chapter 8](#) and [Chapter 9](#)), and in [Chapter 10](#) ([Section 10.2.1](#)) we'll learn how to protect pages from improper access. Finally, in [Chapter 11](#) and [Chapter 12](#) we'll add account activations (thereby confirming a valid email address) and password resets. Taken together, the material in [Chapter 6](#) through [Chapter 12](#) develops a full Rails login and authentication system. As you may know, there are various pre-built authentication solutions for Rails; [Box 6.1](#) explains why, at least at first, it's probably a better idea to roll your own.

Box 6.1. Rolling your own authentication system

Virtually all web applications require a login and authentication system of some sort. As a result, most web frameworks have a plethora of options for implementing such systems, and Rails is no exception. Examples of authentication and authorization systems include [Clearance](#), [Authlogic](#), [Devise](#), and [CanCan](#) (as well as non-Rails-specific solutions built on top of [OpenID](#) or [OAuth](#)). It's reasonable to ask why we should reinvent the wheel. Why not just use an off-the-shelf solution instead of rolling our own?

For one, practical experience shows that authentication on most sites requires extensive customization, and modifying a third-party product is often more work than writing the system from scratch. In addition, off-the-shelf systems can be “black boxes”, with potentially mysterious innards; when you write your own system, you are far more likely to understand it. Moreover, recent additions to Rails ([Section 6.3](#)) make it easy to write a custom authentication system. Finally, if you *do* end up using a third-party system later on, you'll be in a much better position to understand and modify it if you've first built one yourself.

6.1 User model

Although the ultimate goal of the next three chapters is to make a signup page for our site (as mocked up in [Figure 6.1](#)), it would do little good now to accept information for new users: we don't currently have any place to put it. Thus, the first step in signing up users is to make a data structure to capture and store their information.

Figure 6.1: A mockup of the user signup page.

In Rails, the default data structure for a data model is called, naturally enough, a *model* (the M in MVC from [Section 1.3.3](#)). The default Rails solution to the problem of persistence is to use a *database* for long-term data storage, and the default library for interacting with the database is called *Active Record*.¹ Active Record comes with a host of methods for creating, saving, and finding data objects, all without having to use the structured query language (SQL)² used by [relational databases](#). Moreover, Rails has a feature called *migrations* to allow data definitions to be written in pure Ruby, without having to learn an SQL data definition language (DDL). The effect is that Rails insulates you almost entirely from the details of the database. In this book, by using SQLite for development and PostgreSQL (via Heroku) for deployment ([Section 1.5](#)), we have developed this theme even further, to the point where we barely ever have to think about how Rails stores data, even for production applications.

As usual, if you're following along using Git for version control, now would be a good time to make a topic branch for modeling users:

```
$ git checkout -b modeling-users
```

6.1.1 Database migrations

You may recall from [Section 4.4.5](#) that we have already encountered, via a custom-built User class, user objects with name and email attributes. That class served as a useful example, but it lacked the critical property of *persistence*: when we created a User object at the Rails console, it disappeared as soon as we exited. Our goal in this section is to create a model for users that won't disappear quite so easily.

As with the User class in [Section 4.4.5](#), we'll start by modeling a user with two attributes, a name and an email address, the latter of which we'll use as a unique username.³ (We'll add an attribute for passwords in [Section 6.3](#).) In [Listing 4.17](#), we did this with Ruby's `attr_accessor` method:

```
class User
  attr_accessor :name, :email
  .
  .
end
```

end

In contrast, when using Rails to model users we don't need to identify the attributes explicitly. As noted briefly above, to store data Rails uses a relational database by default, which consists of *tables* composed of data *rows*, where each row has *columns* of data attributes. For example, to store users with names and email addresses, we'll create a users table with name and email columns (with each row corresponding to one user). An example of such a table appears in [Figure 6.2](#), corresponding to the data model shown in [Figure 6.3](#). ([Figure 6.3](#) is just a sketch; the full data model appears in [Figure 6.4](#).) By naming the columns name and email, we'll let Active Record figure out the User object attributes for us.

users		
id	name	email
1	Michael Hartl	mhartl@example.com
2	Sterling Archer	archer@example.gov
3	Lana Kane	lane@example.gov
4	Mallory Archer	boss@example.gov

Figure 6.2: A diagram of sample data in a users table.

users	
id	integer
name	string
email	string

Figure 6.3: A sketch of the User data model.

You may recall from [Listing 5.38](#) that we created a Users controller (along with a new action) using the command

```
$ rails generate controller Users new
```

The analogous command for making a model is `generate model`, which we can use to generate a User model with name and email attributes, as shown in [Listing 6.1](#).

Listing 6.1: Generating a User model.

```
$ rails generate model User name:string email:string
      invoke  active_record
      create  db/migrate/20160523010738_create_users.rb
      create  app/models/user.rb
      invoke  test_unit
      create  test/models/user_test.rb
      create  test/fixtures/users.yml
```

(Note that, in contrast to the plural convention for controller names, model names are singular: a *Users* controller, but a *User* model.) By passing the optional parameters `name:string` and `email:string`, we tell Rails about the two attributes we want, along with which types those attributes should be (in this case, `string`). Compare this with including the action names in [Listing 3.6](#) and [Listing 5.38](#).

One of the results of the `generate` command in [Listing 6.1](#) is a new file called a *migration*. Migrations provide a way to alter the structure of the database incrementally, so that our data model can adapt to changing requirements. In the case of the User model, the migration is created automatically by the model generation script; it creates a users table with two columns, name and email, as shown in [Listing 6.2](#). (We'll see starting in [Section 6.2.5](#) how to make a migration from scratch.)

Listing 6.2: Migration for the User model (to create a users table). `db/migrate/[timestamp]_create_users.rb`

```
class CreateUsers < ActiveRecord::Migration[5.0]
  def change
    create_table :users do |t|
      t.string :name
      t.string :email

      t.timestamps
    end
  end
end
```

```
end  
end
```

Note that the name of the migration file is prefixed by a *timestamp* based on when the migration was generated. In the early days of migrations, the filenames were prefixed with incrementing integers, which caused conflicts for collaborating teams if multiple programmers had migrations with the same number. Barring the improbable scenario of migrations generated the same second, using timestamps conveniently avoids such collisions.

The migration itself consists of a change method that determines the change to be made to the database. In the case of [Listing 6.2](#), change uses a Rails method called `create_table` to create a table in the database for storing users. The `create_table` method accepts a block ([Section 4.3.2](#)) with one block variable, in this case called `t` (for “table”). Inside the block, the `create_table` method uses the `t` object to create `name` and `email` columns in the database, both of type `string`.⁴ Here the table name is plural (`users`) even though the model name is singular (`User`), which reflects a linguistic convention followed by Rails: a model represents a single user, whereas a database table consists of many users. The final line in the block, `t.timestamps`, is a special command that creates two *magic columns* called `created_at` and `updated_at`, which are timestamps that automatically record when a given user is created and updated. (We’ll see concrete examples of the magic columns starting in [Section 6.1.3](#).) The full data model represented by the migration in [Listing 6.2](#) is shown in [Figure 6.4](#). (Note the addition of the magic columns, which weren’t present in the sketch shown in [Figure 6.3](#).)

users	
id	integer
name	string
email	string
created_at	datetime
updated_at	datetime

Figure 6.4: The User data model produced by [Listing 6.2](#).

We can run the migration, known as “migrating up”, using the `db:migrate` command as follows:

```
$ rails db:migrate
```

(You may recall that we ran this command in a similar context in [Section 2.2](#).) The first time `db:migrate` is run, it creates a file called `db/development.sqlite3`, which is an [SQLite⁵](#) database. We can see the structure of the database by opening `development.sqlite3` with [DB Browser for SQLite](#). (If you’re using the cloud IDE, you should first download the database file to the local disk, as shown in [Figure 6.5](#).) The result appears in [Figure 6.6](#); compare with the diagram in [Figure 6.4](#). You might note that there’s one column in [Figure 6.6](#) not accounted for in the migration: the `id` column. As noted briefly in [Section 2.2](#), this column is created automatically, and is used by Rails to identify each row uniquely.

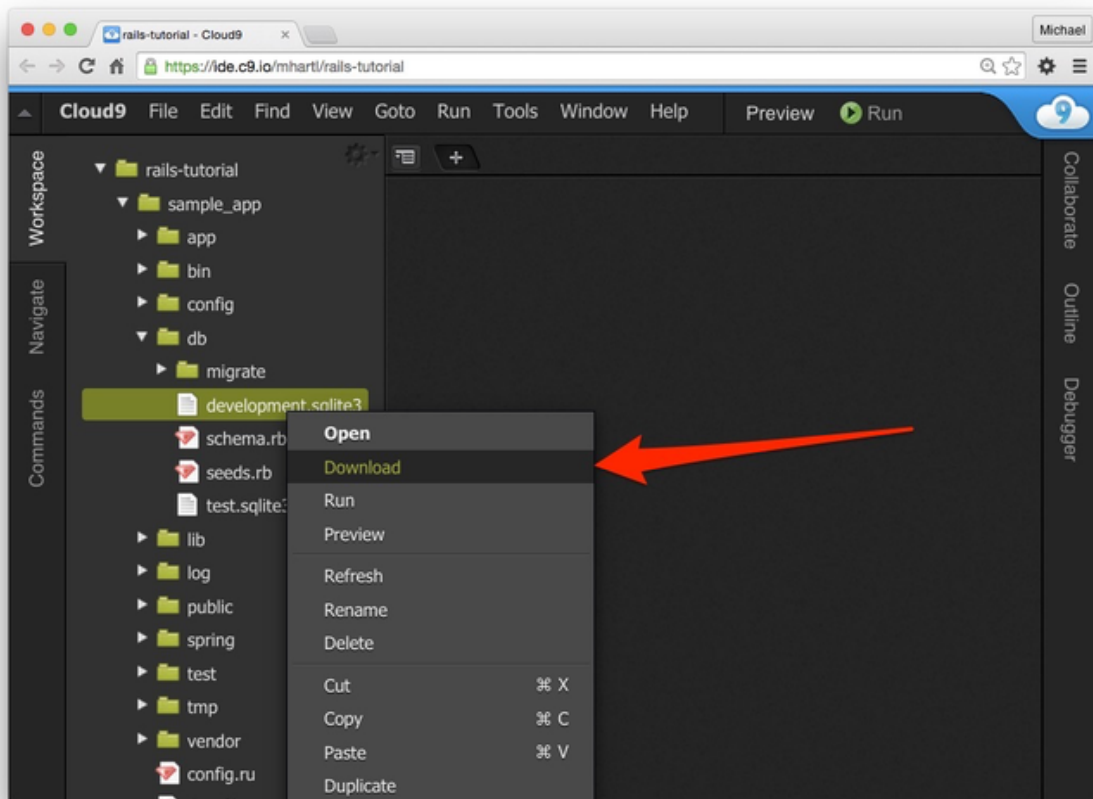


Figure 6.5: Downloading a file from the cloud IDE.

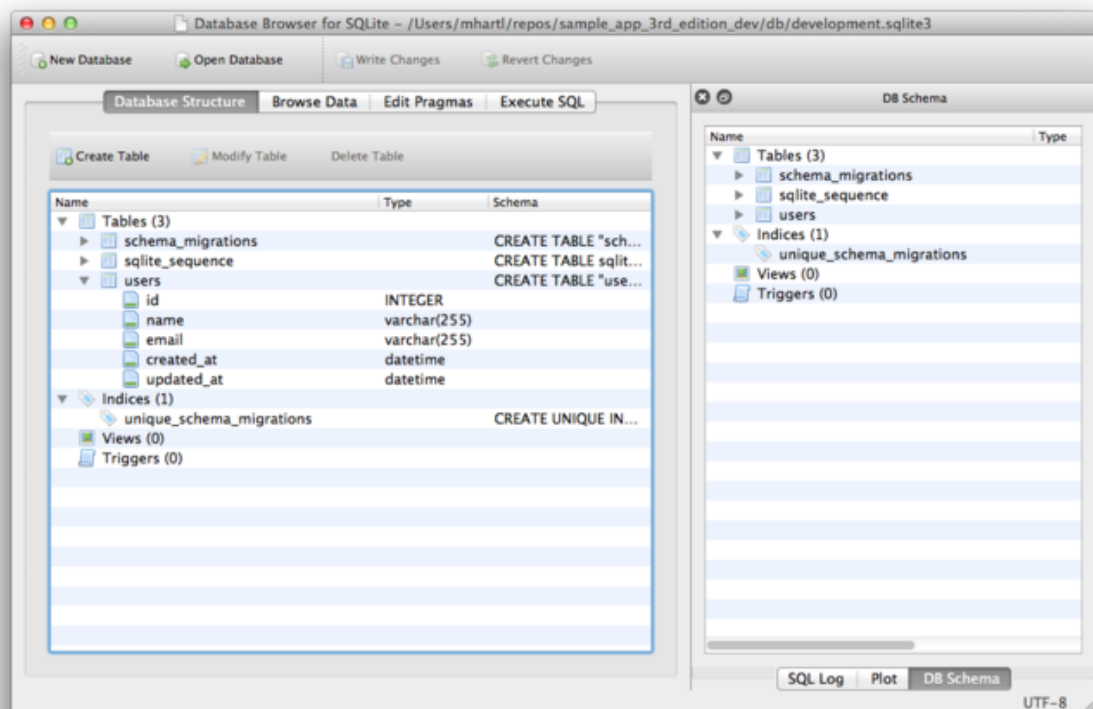


Figure 6.6: The DB Browser for SQLite with our new users table.

Exercises

Solutions to exercises are available for free at railstutorial.org/solutions with any Rails Tutorial purchase. To see other people's answers and to record your own, join the [Learn Enough Society](http://learnenough.com/society) at learnenough.com/society.

1. Rails uses a file called `schema.rb` in the `db/` directory to keep track of the structure of the database (called the *schema*, hence the filename). Examine your local copy of `db/schema.rb` and compare its contents to the migration code in [Listing 6.2](#).
2. Most migrations (including all the ones in this tutorial) are *reversible*, which means we can “migrate down” and undo them with a single command, called `db:rollback`:

```
$ rails db:rollback
```

After running this command, examine `db/schema.rb` to confirm that the rollback was successful. (See [Box 3.1](#) for another technique useful for reversing migrations.) Under the hood, this command executes the `drop_table` command to remove the `users` table from the database. The reason this works is that the `change` method knows that `drop_table` is the inverse of `create_table`, which means that the rollback migration can be easily inferred. In the case of an irreversible migration, such as one to remove a database column, it is necessary to define separate up and down methods in place of the single `change` method. Read about [migrations in the Rails Guides](#) for more information.

3. Re-run the migration by executing `rails db:migrate` again. Confirm that the contents of `db/schema.rb` have been restored.

6.1.2 The model file

We’ve seen how the `User` model generation in [Listing 6.1](#) generated a migration file ([Listing 6.2](#)), and we saw in [Figure 6.6](#) the results of running this migration: it updated a file called `development.sqlite3` by creating a table `users` with columns `id`, `name`, `email`, `created_at`, and `updated_at`. [Listing 6.1](#) also created the model itself. The rest of this section is dedicated to understanding it.

We begin by looking at the code for the `User` model, which lives in the file `user.rb` inside the `app/models/` directory. It is, to put it mildly, very compact ([Listing 6.3](#)).

Listing 6.3: The brand new `User` model. `app/models/user.rb`

```
class User < ApplicationRecord
end
```

Recall from [Section 4.4.2](#) that the syntax `class User < ApplicationRecord` means that the `User` class *inherits* from the `ApplicationRecord` class, which in turn inherits from `ActiveRecord::Base` ([Figure 2.18](#)), so that the `User` model automatically has all the functionality of the `ActiveRecord::Base` class. Of course, this knowledge doesn’t do us any good unless we know what `ActiveRecord::Base` contains, so let’s get started with some concrete examples.

Exercises

Solutions to exercises are available for free at railstutorial.org/solutions with any Rails Tutorial purchase. To see other people’s answers and to record your own, join the Learn Enough Society at learnenough.com/society.

1. In a Rails console, use the technique from [Section 4.4.4](#) to confirm that `User.new` is of class `User` and inherits from `ApplicationRecord`.
2. Confirm that `ApplicationRecord` inherits from `ActiveRecord::Base`.

6.1.3 Creating user objects

As in [Chapter 4](#), our tool of choice for exploring data models is the Rails console. Since we don’t (yet) want to make any changes to our database, we’ll start the console in a *sandbox*:

```
$ rails console --sandbox
Loading development environment in sandbox
Any modifications you make will be rolled back on exit
>>
```

As indicated by the helpful message “Any modifications you make will be rolled back on exit”, when started in a sandbox the console will “roll back” (i.e., undo) any database changes introduced during the session.

In the console session in [Section 4.4.5](#), we created a new user object with `User.new`, which we had access to only after requiring the example user file in [Listing 4.17](#). With models, the situation is different; as you may recall from [Section 4.4.4](#), the Rails console automatically loads the Rails environment, which includes the models. This means that we can make a new user object without any further work:

```
>> User.new
=> #<User id: nil, name: nil, email: nil, created_at: nil, updated_at: nil>
```

We see here the default console representation of a user object.

When called with no arguments, `User.new` returns an object with all `nil` attributes. In [Section 4.4.5](#), we designed the example `User` class to take an *initialization hash* to set the object attributes; that design choice was motivated by Active Record, which allows objects to be initialized in the same way:

```
>> user = User.new(name: "Michael Hartl", email: "mhartl@example.com")
=> #<User id: nil, name: "Michael Hartl", email: "mhartl@example.com",
    created_at: nil, updated_at: nil>
```

Here we see that the name and email attributes have been set as expected.

The notion of *validity* is important for understanding Active Record model objects. We'll explore this subject in more depth in [Section 6.2](#), but for now it's worth noting that our initial user object is valid, which we can verify by calling the boolean `valid?` method on it:

```
>> user.valid?
true
```

So far, we haven't touched the database: `User.new` only creates an object *in memory*, while `user.valid?` merely checks to see if the object is valid. In order to save the `User` object to the database, we need to call the `save` method on the user variable:

```
>> user.save
(0.1ms) SAVEPOINT active_record_1
SQL (0.8ms) INSERT INTO "users" ("name", "email", "created_at",
    "updated_at") VALUES (?, ?, ?, ?) [["name", "Michael Hartl"],
["email", "mhartl@example.com"], ["created_at", 2016-05-23 19:05:58 UTC],
["updated_at", 2016-05-23 19:05:58 UTC]]
(0.1ms) RELEASE SAVEPOINT active_record_1
=> true
```

The `save` method returns `true` if it succeeds and `false` otherwise. (Currently, all saves should succeed because there are as yet no validations; we'll see cases in [Section 6.2](#) when some will fail.) For reference, the Rails console also shows the SQL command corresponding to `user.save` (namely, `INSERT INTO "users"...`). We'll hardly ever need raw SQL in this book,⁶ and I'll omit discussion of the SQL commands from now on, but you can learn a lot by reading the SQL corresponding to Active Record commands.

You may have noticed that the new user object had `nil` values for the `id` and the magic columns `created_at` and `updated_at` attributes. Let's see if our `save` changed anything:

```
>> user
=> #<User id: 1, name: "Michael Hartl", email: "mhartl@example.com",
    created_at: "2016-05-23 19:05:58", updated_at: "2016-05-23 19:05:58">
```

We see that the `id` has been assigned a value of 1, while the magic columns have been assigned the current time and date.⁷ Currently, the `created` and `updated` timestamps are identical; we'll see them differ in [Section 6.1.5](#).

As with the `User` class in [Section 4.4.5](#), instances of the `User` model allow access to their attributes using a dot notation:

```
>> user.name
=> "Michael Hartl"
>> user.email
=> "mhartl@example.com"
>> user.updated_at
=> Mon, 23 May 2016 19:05:58 UTC +00:00
```

As we'll see in [Chapter 7](#), it's often convenient to make and save a model in two steps as we have above, but Active Record also lets you combine them into one step with `User.create`:

```
>> User.create(name: "A Nother", email: "another@example.org")
#<User id: 2, name: "A Nother", email: "another@example.org", created_at:
    "2016-05-23 19:18:46", updated_at: "2016-05-23 19:18:46">
>> foo = User.create(name: "Foo", email: "foo@bar.com")
#<User id: 3, name: "Foo", email: "foo@bar.com", created_at: "2016-05-23
    19:19:06", updated_at: "2016-05-23 19:19:06">
```

Note that `User.create`, rather than returning `true` or `false`, returns the `User` object itself, which we can optionally assign to a variable (such as `foo` in the second command above).

The inverse of create is destroy:

```
>> foo.destroy
(0.1ms) SAVEPOINT active_record_1
SQL (0.2ms) DELETE FROM "users" WHERE "users"."id" = ? [["id", 3]]
(0.1ms) RELEASE SAVEPOINT active_record_1
=> #<User id: 3, name: "Foo", email: "foo@bar.com", created_at: "2016-05-23
19:19:06", updated_at: "2016-05-23 19:19:06">
```

Like create, destroy returns the object in question, though I can't recall ever having used the return value of destroy. In addition, the destroyed object still exists in memory:

```
>> foo
=> #<User id: 3, name: "Foo", email: "foo@bar.com", created_at: "2016-05-23
19:19:06", updated_at: "2016-05-23 19:19:06">
```

So how do we know if we really destroyed an object? And for saved and non-destroyed objects, how can we retrieve users from the database? To answer these questions, we need to learn how to use Active Record to find user objects.

Exercises

Solutions to exercises are available for free at railstutorial.org/solutions with any Rails Tutorial purchase. To see other people's answers and to record your own, join the [Learn Enough Society](http://learnenough.com/society) at learnenough.com/society.

1. Confirm that `user.name` and `user.email` are of class `String`.
2. Of what class are the `created_at` and `updated_at` attributes?

6.1.4 Finding user objects

Active Record provides several options for finding objects. Let's use them to find the first user we created while verifying that the third user (foo) has been destroyed. We'll start with the existing user:

```
>> User.find(1)
=> #<User id: 1, name: "Michael Hartl", email: "mhartl@example.com",
created_at: "2016-05-23 19:05:58", updated_at: "2016-05-23 19:05:58">
```

Here we've passed the id of the user to `User.find`; Active Record returns the user with that id.

Let's see if the user with an id of 3 still exists in the database:

```
>> User.find(3)
ActiveRecord::RecordNotFound: Couldn't find User with ID=3
```

Since we destroyed our third user in [Section 6.1.3](#), Active Record can't find it in the database. Instead, `find` raises an *exception*, which is a way of indicating an exceptional event in the execution of a program—in this case, a nonexistent Active Record id, which causes `find` to raise an `ActiveRecord::RecordNotFound` exception.⁸

In addition to the generic `find`, Active Record also allows us to find users by specific attributes:

```
>> User.find_by(email: "mhartl@example.com")
=> #<User id: 1, name: "Michael Hartl", email: "mhartl@example.com",
created_at: "2016-05-23 19:05:58", updated_at: "2016-05-23 19:05:58">
```

Since we will be using email addresses as usernames, this sort of `find` will be useful when we learn how to let users log in to our site ([Chapter 7](#)). If you're worried that `find_by` will be inefficient if there are a large number of users, you're ahead of the game; we'll cover this issue, and its solution via database indices, in [Section 6.2.5](#).

We'll end with a couple of more general ways of finding users. First, there's `first`:

```
>> User.first
=> #<User id: 1, name: "Michael Hartl", email: "mhartl@example.com",
created_at: "2016-05-23 19:05:58", updated_at: "2016-05-23 19:05:58">
```

Naturally, `first` just returns the first user in the database. There's also `all`:

```
>> User.all
=> #<ActiveRecord::Relation [#<User id: 1, name: "Michael Hartl",
email: "mhartl@example.com", created_at: "2016-05-23 19:05:58",
updated_at: "2016-05-23 19:05:58">, #<User id: 2, name: "A Nother",
```



```
email: "another@example.org", created_at: "2016-05-23 19:18:46",
updated_at: "2016-05-23 19:18:46">]>
```

As you can see from the console output, `User.all` returns all the users in the database as an object of class `ActiveRecord::Relation`, which is effectively an array ([Section 4.3.1](#)).

Exercises

Solutions to exercises are available for free at railstutorial.org/solutions with any Rails Tutorial purchase. To see other people's answers and to record your own, join the [Learn Enough Society](http://learnenough.com/society) at learnenough.com/society.

1. Find the user by name. Confirm that `find_by_name` works as well. (You will often encounter this older style of `find_by` in legacy Rails applications.)
2. For most practical purposes, `User.all` acts like an array, but confirm that in fact it's of class `User::ActiveRecord_Relation`.
3. Confirm that you can find the length of `User.all` by passing it the `length` method ([Section 4.2.3](#)). Ruby's ability to manipulate objects based on how they act rather than on their formal class type is called *duck typing*, based on the aphorism that "If it looks like a duck, and it quacks like a duck, it's probably a duck."

6.1.5 Updating user objects

Once we've created objects, we often want to update them. There are two basic ways to do this. First, we can assign attributes individually, as we did in [Section 4.4.5](#):

```
>> user          # Just a reminder about our user's attributes
=> #<User id: 1, name: "Michael Hartl", email: "mhartl@example.com",
created_at: "2016-05-23 19:05:58", updated_at: "2016-05-23 19:05:58">
>> user.email = "mhartl@example.net"
=> "mhartl@example.net"
>> user.save
=> true
```

Note that the final step is necessary to write the changes to the database. We can see what happens without a `save` by using `reload`, which reloads the object based on the database information:

```
>> user.email
=> "mhartl@example.net"
>> user.email = "foo@bar.com"
=> "foo@bar.com"
>> user.reload.email
=> "mhartl@example.net"
```

Now that we've updated the user by running `user.save`, the magic columns differ, as promised in [Section 6.1.3](#):

```
>> user.created_at
=> "2016-05-23 19:05:58"
>> user.updated_at
=> "2016-05-23 19:08:23"
```

The second main way to update multiple attributes is to use `update_attributes`.⁹

```
>> user.update_attributes(name: "The Dude", email: "dude@abides.org")
=> true
>> user.name
=> "The Dude"
>> user.email
=> "dude@abides.org"
```

The `update_attributes` method accepts a hash of attributes, and on success performs both the update and the save in one step (returning `true` to indicate that the save went through). Note that if any of the validations fail, such as when a password is required to save a record (as implemented in [Section 6.3](#)), the call to `update_attributes` will fail. If we need to update only a single attribute, using the singular `update_attribute` bypasses this restriction by skipping the validations:

```
>> user.update_attribute(:name, "El Duderino")
=> true
>> user.name
=> "El Duderino"
```

Exercises

Solutions to exercises are available for free at railstutorial.org/solutions with any Rails Tutorial purchase. To see other people's answers and to record your own, join the [Learn Enough Society](http://learnenough.com/society) at learnenough.com/society.

1. Update the user's name using assignment and a call to `save`.
2. Update the user's email address using a call to `update_attributes`.
3. Confirm that you can change the magic columns directly by updating the `created_at` column using assignment and a `save`. Use the value `1.year.ago`, which is a Rails way to create a timestamp one year before the present time.

6.2 User validations

The User model we created in [Section 6.1](#) now has working name and email attributes, but they are completely generic: any string (including an empty one) is currently valid in either case. And yet, names and email addresses are more specific than this. For example, name should be non-blank, and email should match the specific format characteristic of email addresses. Moreover, since we'll be using email addresses as unique usernames when users log in, we shouldn't allow email duplicates in the database.

In short, we shouldn't allow name and email to be just any strings; we should enforce certain constraints on their values. Active Record allows us to impose such constraints using *validations* (seen briefly before in [Section 2.3.2](#)). In this section, we'll cover several of the most common cases, validating *presence*, *length*, *format* and *uniqueness*. In [Section 6.3.2](#) we'll add a final common validation, *confirmation*. And we'll see in [Section 7.3](#) how validations give us convenient error messages when users make submissions that violate them.

6.2.1 A validity test

As noted in [Box 3.3](#), test-driven development isn't always the right tool for the job, but model validations are exactly the kind of features for which TDD is a perfect fit. It's difficult to be confident that a given validation is doing exactly what we expect it to without writing a failing test and then getting it to pass.

Our method will be to start with a *valid* model object, set one of its attributes to something we want to be invalid, and then test that it in fact is invalid. As a safety net, we'll first write a test to make sure the initial model object is valid. This way, when the validation tests fail we'll know it's for the right reason (and not because the initial object was invalid in the first place).

To get us started, the command in [Listing 6.1](#) produced an initial test for testing users, though in this case it's practically blank ([Listing 6.4](#)).

Listing 6.4: The practically blank default User test. `test/models/user_test.rb`

```
require 'test_helper'

class UserTest < ActiveSupport::TestCase
  # test "the truth" do
  #   assert true
  # end
end
```

To write a test for a valid object, we'll create an initially valid User model object `@user` using the special setup method (discussed briefly in the [Chapter 3](#) exercises), which automatically gets run before each test. Because `@user` is an instance variable, it's automatically available in all the tests, and we can test its validity using the `valid?` method ([Section 6.1.3](#)). The result appears in [Listing 6.5](#).

Listing 6.5: A test for an initially valid user. **green** `test/models/user_test.rb`

```
require 'test_helper'

class UserTest < ActiveSupport::TestCase

  def setup
    @user = User.new(name: "Example User", email: "user@example.com")
  end

  test "should be valid" do
    assert @user.valid?
  end
end
```

[Listing 6.5](#) uses the plain `assert` method, which in this case succeeds if `@user.valid?` returns `true` and fails if it returns `false`.

Because our `User` model doesn't currently have any validations, the initial test should pass:

Listing 6.6: **green**

```
$ rails test:models
```

Here we've used `rails test:models` to run just the model tests (compare to `rails test:integration` from [Section 5.3.4](#)).

Exercises

Solutions to exercises are available for free at railstutorial.org/solutions with any Rails Tutorial purchase. To see other people's answers and to record your own, join the Learn Enough Society at learnenough.com/society.

1. In the console, confirm that a new user is currently valid.
2. Confirm that the user created in [Section 6.1.3](#) is also valid.

6.2.2 Validating presence

Perhaps the most elementary validation is *presence*, which simply verifies that a given attribute is present. For example, in this section we'll ensure that both the name and email fields are present before a user gets saved to the database. In [Section 7.3.3](#), we'll see how to propagate this requirement up to the signup form for creating new users.

We'll start with a test for the presence of a name attribute by building on the test in [Listing 6.5](#). As seen in [Listing 6.7](#), all we need to do is set the `@user` variable's name attribute to a blank string (in this case, a string of spaces) and then check (using the `assert_not` method) that the resulting `User` object is not valid.

Listing 6.7: A test for validation of the name attribute. **red** `test/models/user_test.rb`

```
require 'test_helper'

class UserTest < ActiveSupport::TestCase
  def setup
    @user = User.new(name: "Example User", email: "user@example.com")
  end

  test "should be valid" do
    assert @user.valid?
  end

  test "name should be present" do
    @user.name = " "
    assert_not @user.valid?
  end
end
```

At this point, the model tests should be **red**:

Listing 6.8: **red**

```
$ rails test:models
```

As we saw briefly before in the [Chapter 2](#) exercises, the way to validate the presence of the name attribute is to use the `validates` method with argument `presence: true`, as shown in [Listing 6.9](#). The `presence: true` argument is a one-element *options hash*; recall from [Section 4.3.4](#) that curly braces are optional when passing hashes as the final argument in a method. (As noted in [Section 5.1.1](#), the use of options hashes is a recurring theme in Rails.)

Listing 6.9: Validating the presence of a name attribute. **green** `app/models/user.rb`

```
class User < ApplicationRecord
  validates :name, presence: true
end
```

[Listing 6.9](#) may look like magic, but `validates` is just a method. An equivalent formulation of [Listing 6.9](#) using parentheses is as follows:

```
class User < ApplicationRecord
  validates(:name, presence: true)
end
```

Let's drop into the console to see the effects of adding a validation to our User model:¹⁰

```
$ rails console --sandbox
>> user = User.new(name: "", email: "mhartl@example.com")
>> user.valid?
=> false
```

Here we check the validity of the user variable using the `valid?` method, which returns `false` when the object fails one or more validations, and `true` when all validations pass. In this case, we only have one validation, so we know which one failed, but it can still be helpful to check using the errors object generated on failure:

```
>> user.errors.full_messages
=> ["Name can't be blank"]
```

(The error message is a hint that Rails validates the presence of an attribute using the `blank?` method, which we saw at the end of [Section 4.4.3](#).)

Because the user isn't valid, an attempt to save the user to the database automatically fails:

```
>> user.save
=> false
```

As a result, the test in [Listing 6.7](#) should now be **green**:

Listing 6.10: **green**

```
$ rails test:models
```

Following the model in [Listing 6.7](#), writing a test for email attribute presence is easy ([Listing 6.11](#)), as is the application code to get it to pass ([Listing 6.12](#)).

Listing 6.11: A test for validation of the email attribute. **red** test/models/user_test.rb

```
require 'test_helper'

class UserTest < ActiveSupport::TestCase

  def setup
    @user = User.new(name: "Example User", email: "user@example.com")
  end

  test "should be valid" do
    assert @user.valid?
  end

  test "name should be present" do
    @user.name = ""
    assert_not @user.valid?
  end

  test "email should be present" do
    @user.email = ""
    assert_not @user.valid?
  end
end
```

Listing 6.12: Validating the presence of an email attribute. **green** app/models/user.rb

```
class User < ApplicationRecord
  validates :name, presence: true
  validates :email, presence: true
end
```

At this point, the presence validations are complete, and the test suite should be **green**:

Listing 6.13: **green**

```
$ rails test
```

Exercises

Solutions to exercises are available for free at railstutorial.org/solutions with any Rails Tutorial purchase. To see other people's answers and to record your own, join the [Learn Enough Society](http://learnenough.com/society) at learnenough.com/society.

1. Make a new user called u and confirm that it's initially invalid. What are the full error messages?
2. Confirm that `u.errors.messages` is a hash of errors. How would you access just the email errors?

6.2.3 Length validation

We've constrained our User model to require a name for each user, but we should go further: the user's names will be displayed on the sample site, so we should enforce some limit on their length. With all the work we did in [Section 6.2.2](#), this step is easy.

There's no science to picking a maximum length; we'll just pull 50 out of thin air as a reasonable upper bound, which means verifying that names of 51 characters are too long. In addition, although it's unlikely ever to be a problem, there's a chance that a user's email address could overrun the maximum length of strings, which for many databases is 255. Because the format validation in [Section 6.2.4](#) won't enforce such a constraint, we'll add one in this section for completeness. [Listing 6.14](#) shows the resulting tests.

Listing 6.14: A test for name length validation. **red** `test/models/user_test.rb`

```
require 'test_helper'

class UserTest < ActiveSupport::TestCase

  def setup
    @user = User.new(name: "Example User", email: "user@example.com")
  end
  .
  .
  .
  test "name should not be too long" do
    @user.name = "a" * 51
    assert_not @user.valid?
  end

  test "email should not be too long" do
    @user.email = "a" * 244 + "@example.com"
    assert_not @user.valid?
  end
end
```

For convenience, we've used "string multiplication" in [Listing 6.14](#) to make a string 51 characters long. We can see how this works using the console:

```
>> "a" * 51
=> "aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa"
>> ("a" * 51).length
=> 51
```

The email length validation arranges to make a valid email address that's one character too long:

```
>> "a" * 244 + "@example.com"
=> "aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa
aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa
aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa
aaaaaaaaaaaaa@example.com"
>> ("a" * 244 + "@example.com").length
=> 256
```

At this point, the tests in [Listing 6.14](#) should be **red**:

Listing 6.15: **red**

```
$ rails test
```

To get them to pass, we need to use the validation argument to constrain length, which is just `length`, along with the `maximum` parameter to enforce the upper bound ([Listing 6.16](#)).

Listing 6.16: Adding a length validation for the name attribute. **green** app/models/user.rb

```
class User < ApplicationRecord
  validates :name, presence: true, length: { maximum: 50 }
  validates :email, presence: true, length: { maximum: 255 }
end
```

Now the tests should be **green**:

Listing 6.17: **green**

```
$ rails test
```

With our test suite passing again, we can move on to a more challenging validation: email format.

Exercises

Solutions to exercises are available for free at railstutorial.org/solutions with any Rails Tutorial purchase. To see other people's answers and to record your own, join the [Learn Enough Society](http://learnenough.com/society) at learnenough.com/society.

1. Make a new user with too-long name and email and confirm that it's not valid.
2. What are the error messages generated by the length validation?

6.2.4 Format validation

Our validations for the name attribute enforce only minimal constraints—any non-blank name under 51 characters will do—but of course the email attribute must satisfy the more stringent requirement of being a valid email address. So far we've only rejected blank email addresses; in this section, we'll require email addresses to conform to the familiar pattern `user@example.com`.

Neither the tests nor the validation will be exhaustive, just good enough to accept most valid email addresses and reject most invalid ones. We'll start with a couple of tests involving collections of valid and invalid addresses. To make these collections, it's worth knowing about the useful `%w[]` technique for making arrays of strings, as seen in this console session:

```
>> %w[foo bar baz]
=> ["foo", "bar", "baz"]
>> addresses = %w[USER@foo.COM THE_US-ER@foo.bar.org first.last@foo.jp]
=> ["USER@foo.COM", "THE_US-ER@foo.bar.org", "first.last@foo.jp"]
>> addresses.each do |address|
  ?> puts address
  >> end
USER@foo.COM
THE_US-ER@foo.bar.org
first.last@foo.jp
```

Here we've iterated over the elements of the `addresses` array using the `each` method (Section 4.3.2). With this technique in hand, we're ready to write some basic email format validation tests.

Because email format validation is tricky and error-prone, we'll start with some passing tests for *valid* email addresses to catch any errors in the validation. In other words, we want to make sure not just that invalid email addresses like `user@example.com` are rejected, but also that valid addresses like `user@example.com` are accepted, even after we impose the validation constraint. (Right now, of course, they'll be accepted because all non-blank email addresses are currently valid.) The result for a representative sample of valid email addresses appears in Listing 6.18.

Listing 6.18: Tests for valid email formats. **green** test/models/user_test.rb

```
require 'test_helper'

class UserTest < ActiveSupport::TestCase

  def setup
    @user = User.new(name: "Example User", email: "user@example.com")
  end

  .
  .
  .

  test "email validation should accept valid addresses" do
    valid_addresses = %w[user@example.com USER@foo.COM A_US-ER@foo.bar.org
```

```

        first.last@foo.jp alice+bob@baz.cn]
    valid_addresses.each do |valid_address|
      @user.email = valid_address
      assert @user.valid?, "#{valid_address.inspect} should be valid"
    end
  end
end

```

Note that we've included an optional second argument to the assertion with a custom error message, which in this case identifies the address causing the test to fail:

```
assert @user.valid?, "#{valid_address.inspect} should be valid"
```

(This uses the interpolated `inspect` method mentioned in [Section 4.3.3](#).) Including the specific address that causes any failure is especially useful in a test with an `each` loop like [Listing 6.18](#); otherwise, any failure would merely identify the line number, which is the same for all the email addresses, and which wouldn't be sufficient to identify the source of the problem.

Next we'll add tests for the *invalidity* of a variety of invalid email addresses, such as `user@example.com` (comma in place of dot) and `user_at_foo.org` (missing the '@' sign). As in [Listing 6.18](#), [Listing 6.19](#) includes a custom error message to identify the exact address causing any failure.

Listing 6.19: Tests for email format validation. **red** `test/models/user_test.rb`

```

require 'test_helper'

class UserTest < ActiveSupport::TestCase
  def setup
    @user = User.new(name: "Example User", email: "user@example.com")
  end
  .
  .
  .
  test "email validation should reject invalid addresses" do
    invalid_addresses = %w[user@example.com user_at_foo.org user.name@example.
                          foo@bar_baz.com foo@bar+baz.com]
    invalid_addresses.each do |invalid_address|
      @user.email = invalid_address
      assert_not @user.valid?, "#{invalid_address.inspect} should be invalid"
    end
  end
end

```

At this point, the tests should be **red**:

Listing 6.20: **red**

```
$ rails test
```

The application code for email format validation uses the `format` validation, which works like this:

```
validates :email, format: { with: /<regular expression>/ }
```

This validates the attribute with the given *regular expression* (or *regex*), which is a powerful (and often cryptic) language for matching patterns in strings. This means we need to construct a regular expression to match valid email addresses while *not* matching invalid ones.

There actually exists a full regex for matching email addresses according to the official email standard, but it's enormous, obscure, and quite possibly counter-productive.¹¹ In this tutorial, we'll adopt a more pragmatic regex that has proven to be robust in practice. Here's what it looks like:

```
VALID_EMAIL_REGEX = /\A[\w+\-\.]+\@[a-z\d\-\.\+\.][a-z]+\z/i
```

To help understand where this comes from, [Table 6.1](#) breaks it into bite-sized pieces.¹²

Expression	Meaning
<code>/\A[\w+\-\.]+\@[a-z\d\-\.\+\.][a-z]+\z/i</code>	full regex
<code>/</code>	start of regex

<code>\A</code>	match start of a string
<code>[\w+\-\.]+</code>	at least one word character, plus, hyphen, or dot
<code>@</code>	literal “at sign”
<code>[a-z\d\-\-\.]+</code>	at least one letter, digit, hyphen, or dot
<code>\.</code>	literal dot
<code>[a-z]+</code>	at least one letter
<code>\z</code>	match end of a string
<code>/</code>	end of regex
<code>i</code>	case-insensitive

Table 6.1: Breaking down the valid email regex.

Although you can learn a lot by studying [Table 6.1](#), to really understand regular expressions I consider using an interactive regular expression matcher like [Rubular](#) to be essential ([Figure 6.7](#)).¹³ The Rubular website has a beautiful interactive interface for making regular expressions, along with a handy regex quick reference. I encourage you to study [Table 6.1](#) with a browser window open to Rubular—no amount of reading about regular expressions can replace playing with them interactively. (Note: If you use the regex from [Table 6.1](#) in Rubular, I recommend leaving off the `\A` and `\z` characters so that you can match more than one email address at a time in the given test string. Also note that the regex consists of the characters *inside* the slashes `/ . . /`, so you should omit those when using Rubular.)

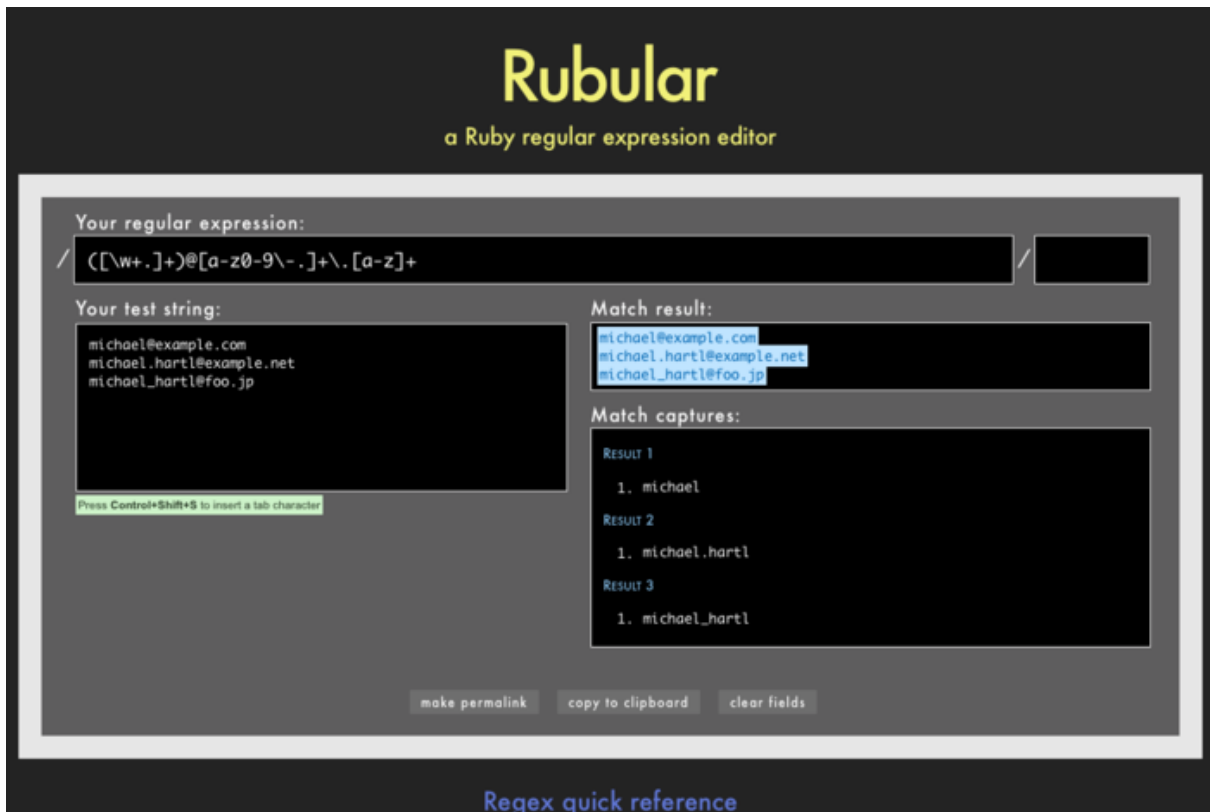


Figure 6.7: The awesome [Rubular](#) regular expression editor.

Applying the regular expression from [Table 6.1](#) to the email format validation yields the code in [Listing 6.21](#).

Listing 6.21: Validating the email format with a regular expression. [green](#) `app/models/user.rb`

```
class User < ApplicationRecord
  validates :name, presence: true, length: { maximum: 50 }
  VALID_EMAIL_REGEX = /\A[\w+\-\. ]+([a-z\d\-\-\. ]+\.[a-z ]+)\z/i
  validates :email, presence: true, length: { maximum: 255 },
    format: { with: VALID_EMAIL_REGEX }
end
```

Here the regex `VALID_EMAIL_REGEX` is a *constant*, indicated in Ruby by a name starting with a capital letter. The code

```
VALID_EMAIL_REGEX = /\A[\w+\-\. ]+([a-z\d\-\-\. ]+\.[a-z ]+)\z/i
validates :email, presence: true, length: { maximum: 255 },
```

```
format: { with: VALID_EMAIL_REGEX }
```

ensures that only email addresses that match the pattern will be considered valid. (The expression above has one minor weakness: it allows invalid addresses that contain consecutive dots, such as `foo@bar..com`. Updating the regex in [Listing 6.21](#) to fix this blemish is left as an exercise ([Section 6.2.4.1](#)).

At this point, the tests should be **green**:

Listing 6.22: **green**

```
$ rails test:models
```

This means that there's only one constraint left: enforcing email uniqueness.

Exercises

Solutions to exercises are available for free at railstutorial.org/solutions with any Rails Tutorial purchase. To see other people's answers and to record your own, join the [Learn Enough Society](http://learnenough.com/society) at learnenough.com/society.

1. By pasting in the valid addresses from [Listing 6.18](#) and invalid addresses from [Listing 6.19](#) into the test string area at Rubular, confirm that the regex from [Listing 6.21](#) matches all of the valid addresses and none of the invalid ones.
2. As noted above, the email regex in [Listing 6.21](#) allows invalid email addresses with consecutive dots in the domain name, i.e., addresses of the form `foo@bar..com`. Add this address to the list of invalid addresses in [Listing 6.19](#) to get a failing test, and then use the more complicated regex shown in [Listing 6.23](#) to get the test to pass.
3. Add `foo@bar..com` to the list of addresses at Rubular, and confirm that the regex shown in [Listing 6.23](#) matches all the valid addresses and none of the invalid ones.

Listing 6.23: Disallowing double dots in email domain names. **green** `app/models/user.rb`

```
class User < ApplicationRecord
  validates :name, presence: true, length: { maximum: 50 }
  VALID_EMAIL_REGEX = /\A[\w+\-\.]+\@[a-z\d\-\-]+\([\.][a-z\d\-\-]+\)*\.[a-z]+\z/i
  validates :email, presence: true, length: { maximum: 255 },
    format: { with: VALID_EMAIL_REGEX }
end
```

6.2.5 Uniqueness validation

To enforce uniqueness of email addresses (so that we can use them as usernames), we'll be using the `:unique` option to the `validates` method. But be warned: there's a *major* caveat, so don't just skim this section—read it carefully.

We'll start with some short tests. In our previous model tests, we've mainly used `User.new`, which just creates a Ruby object in memory, but for uniqueness tests we actually need to put a record into the database.¹⁴ The initial duplicate email test appears in [Listing 6.24](#).

Listing 6.24: A test for the rejection of duplicate email addresses. **red** `test/models/user_test.rb`

```
require 'test_helper'

class UserTest < ActiveSupport::TestCase

  def setup
    @user = User.new(name: "Example User", email: "user@example.com")
  end
  .
  .
  .
  test "email addresses should be unique" do
    duplicate_user = @user.dup
    @user.save
    assert_not duplicate_user.valid?
  end
end
```

The method here is to make a user with the same email address as `@user` using `@user.dup`, which creates a duplicate user with the same attributes. Since we then save `@user`, the duplicate user has an email address that

already exists in the database, and hence should not be valid.

We can get the new test in [Listing 6.24](#) to pass by adding uniqueness: true to the email validation, as shown in [Listing 6.25](#).

Listing 6.25: Validating the uniqueness of email addresses. **green** app/models/user.rb

```
class User < ApplicationRecord
  validates :name, presence: true, length: { maximum: 50 }
  VALID_EMAIL_REGEX = /\A[\w+\-\.]+\@[a-z\d\-\.]+\.[a-z]+\z/i
  validates :email, presence: true, length: { maximum: 255 },
    format: { with: VALID_EMAIL_REGEX },
    uniqueness: true
end
```

We're not quite done, though. Email addresses are typically processed as if they were case-insensitive—i.e., foo@bar.com is treated the same as F00@BAR.COM or Fo0@BAr.coM—so our validation should incorporate this as well.¹⁵ It's thus important to test for case-insensitivity, which we do with the code in [Listing 6.26](#).

Listing 6.26: Testing case-insensitive email uniqueness. **red** test/models/user_test.rb

```
require 'test_helper'

class UserTest < ActiveSupport::TestCase

  def setup
    @user = User.new(name: "Example User", email: "user@example.com")
  end
  .
  .
  .
  test "email addresses should be unique" do
    duplicate_user = @user.dup
    duplicate_user.email = @user.email.upcase
    @user.save
    assert_not duplicate_user.valid?
  end
end
```

Here we are using the upcase method on strings (seen briefly in [Section 4.3.2](#)). This test does the same thing as the initial duplicate email test, but with an upper-case email address instead. If this test feels a little abstract, go ahead and fire up the console:

```
$ rails console --sandbox
>> user = User.create(name: "Example User", email: "user@example.com")
>> user.email.upcase
=> "USER@EXAMPLE.COM"
>> duplicate_user = user.dup
>> duplicate_user.email = user.email.upcase
>> duplicate_user.valid?
=> true
```

Of course, duplicate_user.valid? is currently true because the uniqueness validation is case-sensitive, but we want it to be false. Fortunately, :uniqueness accepts an option, :case_sensitive, for just this purpose ([Listing 6.27](#)).

Listing 6.27: Validating the uniqueness of email addresses, ignoring case. **green** app/models/user.rb

```
class User < ApplicationRecord
  validates :name, presence: true, length: { maximum: 50 }
  VALID_EMAIL_REGEX = /\A[\w+\-\.]+\@[a-z\d\-\.]+\.[a-z]+\z/i
  validates :email, presence: true, length: { maximum: 255 },
    format: { with: VALID_EMAIL_REGEX },
    uniqueness: { case_sensitive: false }
end
```

Note that we have simply replaced true in [Listing 6.25](#) with case_sensitive: false in [Listing 6.27](#). (Rails infers that uniqueness should be true as well.)

At this point, our application—with an important caveat—enforces email uniqueness, and our test suite should pass:

Listing 6.28: green

```
$ rails test
```

There's just one small problem, which is that *the Active Record uniqueness validation does not guarantee uniqueness at the database level*. Here's a scenario that explains why:

1. Alice signs up for the sample app, with address `alice@wonderland.com`.
2. Alice accidentally clicks on "Submit" *twice*, sending two requests in quick succession.
3. The following sequence occurs: request 1 creates a user in memory that passes validation, request 2 does the same, request 1's user gets saved, request 2's user gets saved.
4. Result: two user records with the exact same email address, despite the uniqueness validation

If the above sequence seems implausible, believe me, it isn't: it can happen on any Rails website with significant traffic (which I once learned the hard way). Luckily, the solution is straightforward to implement: we just need to enforce uniqueness at the database level as well as at the model level. Our method is to create a database *index* on the email column (Box 6.2), and then require that the index be unique.

Box 6.2. Database indices

When creating a column in a database, it is important to consider whether we will need to *find* records by that column. Consider, for example, the email attribute created by the migration in Listing 6.2. When we allow users to log in to the sample app starting in Chapter 7, we will need to find the user record corresponding to the submitted email address. Unfortunately, based on the naïve data model, the only way to find a user by email address is to look through *each* user row in the database and compare its email attribute to the given email—which means we might have to examine *every* row (since the user could be the last one in the database). This is known in the database business as a *full-table scan*, and for a real site with thousands of users it is a *Bad Thing*.

Putting an index on the email column fixes the problem. To understand a database index, it's helpful to consider the analogy of a book index. In a book, to find all the occurrences of a given string, say "foobar", you would have to scan each page for "foobar"—the paper version of a full-table scan. With a book index, on the other hand, you can just look up "foobar" in the index to see all the pages containing "foobar". A database index works essentially the same way.

The email index represents an update to our data modeling requirements, which (as discussed in Section 6.1.1) is handled in Rails using migrations. We saw in Section 6.1.1 that generating the User model automatically created a new migration (Listing 6.2); in the present case, we are adding structure to an existing model, so we need to create a migration directly using the migration generator:

```
$ rails generate migration add_index_to_users_email
```

Unlike the migration for users, the email uniqueness migration is not pre-defined, so we need to fill in its contents with Listing 6.29.¹⁶

Listing 6.29: The migration for enforcing email uniqueness.
`db/migrate/[timestamp]_add_index_to_users_email.rb`

```
class AddIndexToUsersEmail < ActiveRecord::Migration[5.0]
  def change
    add_index :users, :email, unique: true
  end
end
```

This uses a Rails method called `add_index` to add an index on the email column of the users table. The index by itself doesn't enforce uniqueness, but the option `unique: true` does.

The final step is to migrate the database:

```
$ rails db:migrate
```

(If this fails, try exiting any running sandbox console sessions, which can lock the database and prevent migrations.)

At this point, the test suite should be **red** due to a violation of the uniqueness constraint in the *fixtures*, which contain sample data for the test database. User fixtures were generated automatically in Listing 6.1, and as shown in Listing 6.30 the email addresses are not unique. (They're not *valid* either, but fixture data doesn't get run through the validations.)

Listing 6.30: The default user fixtures. **red** `test/fixtures/users.yml`

```
# Read about fixtures at http://api.rubyonrails.org/classes/ActiveRecord/
# FixtureSet.html
```

```
one:
  name: MyString
  email: MyString

two:
  name: MyString
  email: MyString
```

Because we won't need fixtures until [Chapter 8](#), for now we'll just remove them, leaving an empty fixtures file ([Listing 6.31](#)).

Listing 6.31: An empty fixtures file. **green** test/fixtures/users.yml

```
# empty
```

Having addressed the uniqueness caveat, there's one more change we need to make to be assured of email uniqueness. Some database adapters use case-sensitive indices, considering the strings "Foo@ExAMPlE.CoM" and "foo@example.com" to be distinct, but our application treats those addresses as the same. To avoid this incompatibility, we'll standardize on all lower-case addresses, converting "Foo@ExAMPlE.CoM" to "foo@example.com" before saving it to the database. The way to do this is with a [callback](#), which is a method that gets invoked at a particular point in the lifecycle of an Active Record object. In the present case, that point is before the object is saved, so we'll use a `before_save` callback to downcase the email attribute before saving the user.¹⁷ The result appears in [Listing 6.32](#). (This is just a first implementation; we'll discuss this subject again in [Section 11.1](#), where we'll use the preferred *method reference* convention for defining callbacks.)

Listing 6.32: Ensuring email uniqueness by downcasing the email attribute. **green** app/models/user.rb

```
class User < ApplicationRecord
  before_save { self.email = email.downcase }
  validates :name, presence: true, length: { maximum: 50 }
  VALID_EMAIL_REGEX = /\A[\w+\-\.]+\@[a-z\d\-\.\.]+\.[a-z]+\z/i
  validates :email, presence: true, length: { maximum: 255 },
    format: { with: VALID_EMAIL_REGEX },
    uniqueness: { case_sensitive: false }
end
```

The code in [Listing 6.32](#) passes a block to the `before_save` callback and sets the user's email address to a lower-case version of its current value using the `downcase` string method. (Writing a test for email downcasing is left as an exercise ([Section 6.2.5.1](#)).)

In [Listing 6.32](#), we could have written the assignment as

```
self.email = self.email.downcase
```

(where `self` refers to the current user), but inside the `User` model the `self` keyword is optional on the right-hand side:

```
self.email = email.downcase
```

We encountered this idea briefly in the context of `reverse` in the `palindrome` method ([Section 4.4.2](#)), which also noted that `self` is *not* optional in an assignment, so

```
email = email.downcase
```

wouldn't work. (We'll discuss this subject in more depth in [Section 9.1](#).)

At this point, the Alice scenario above will work fine: the database will save a user record based on the first request, and it will reject the second save because the duplicate email address violates the uniqueness constraint. (An error will appear in the Rails log, but that doesn't do any harm.) Moreover, adding this index on the email attribute accomplishes a second goal, alluded to briefly in [Section 6.1.4](#): as noted in [Box 6.2](#), the index on the email attribute fixes a potential efficiency problem by preventing a full-table scan when finding users by email address.

Exercises

Solutions to exercises are available for free at railstutorial.org/solutions with any Rails Tutorial purchase. To see other people's answers and to record your own, join the [Learn Enough Society](http://learnenough.com/society) at learnenough.com/society.

1. Add a test for the email downcasing from [Listing 6.32](#), as shown in [Listing 6.33](#). This test uses the `reload` method for reloading a value from the database and the `assert_equal` method for testing equality. To verify that [Listing 6.33](#) tests the right thing, comment out the `before_save` line to get to **red**, then uncomment it to get to **green**.
2. By running the test suite, verify that the `before_save` callback can be written using the “bang” method `email.downcase!` to modify the email attribute directly, as shown in [Listing 6.34](#).

Listing 6.33: A test for the email downcasing from [Listing 6.32](#). `test/models/user_test.rb`

```
require 'test_helper'

class UserTest < ActiveSupport::TestCase
  def setup
    @user = User.new(name: "Example User", email: "user@example.com")
  end
  .
  .
  .
  test "email addresses should be unique" do
    duplicate_user = @user.dup
    duplicate_user.email = @user.email.upcase
    @user.save
    assert_not duplicate_user.valid?
  end

  test "email addresses should be saved as lower-case" do
    mixed_case_email = "Foo@ExAMPlE.CoM"
    @user.email = mixed_case_email
    @user.save
    assert_equal mixed_case_email.downcase, @user.reload.email
  end
end
```

Listing 6.34: An alternate callback implementation. **green** `app/models/user.rb`

```
class User < ApplicationRecord
  before_save { email.downcase! }
  validates :name, presence: true, length: { maximum: 50 }
  VALID_EMAIL_REGEX = /\A[\w+\-\.]+\@[a-z\d\-\.\+]\.[a-z]+\z/i
  validates :email, presence: true, length: { maximum: 255 },
    format: { with: VALID_EMAIL_REGEX },
    uniqueness: { case_sensitive: false }
end
```

6.3 Adding a secure password

Now that we’ve defined validations for the name and email fields, we’re ready to add the last of the basic User attributes: a secure password. The method is to require each user to have a password (with a password confirmation), and then store a *hashed* version of the password in the database. (There is some potential for confusion here. In the present context, a *hash* refers not to the Ruby data structure from [Section 4.3.3](#) but rather to the result of applying an irreversible [hash function](#) to input data.) We’ll also add a way to *authenticate* a user based on a given password, a method we’ll use in [Chapter 8](#) to allow users to log in to the site.

The method for authenticating users will be to take a submitted password, hash it, and compare the result to the hashed value stored in the database. If the two match, then the submitted password is correct and the user is authenticated. By comparing hashed values instead of raw passwords, we will be able to authenticate users without storing the passwords themselves. This means that, even if our database is compromised, our users’ passwords will still be secure.

6.3.1 A hashed password

Most of the secure password machinery will be implemented using a single Rails method called `has_secure_password`, which we’ll include in the User model as follows:

```
class User < ApplicationRecord
  .
  .
  .
```

```
has_secure_password
end
```

When included in a model as above, this one method adds the following functionality:

- The ability to save a securely hashed password_digest attribute to the database
- A pair of virtual attributes¹⁸ (password and password_confirmation), including presence validations upon object creation and a validation requiring that they match
- An authenticate method that returns the user when the password is correct (and false otherwise)

The only requirement for has_secure_password to work its magic is for the corresponding model to have an attribute called password_digest. (The name *digest* comes from the terminology of [cryptographic hash functions](#). In this context, *hashed password* and *password digest* are synonyms.)¹⁹ In the case of the User model, this leads to the data model shown in [Figure 6.8](#).

users	
id	integer
name	string
email	string
created_at	datetime
updated_at	datetime
password_digest	string

Figure 6.8: The User data model with an added password_digest attribute.

To implement the data model in [Figure 6.8](#), we first generate an appropriate migration for the password_digest column. We can choose any migration name we want, but it's convenient to end the name with to_users, since in this case Rails automatically constructs a migration to add columns to the users table. The result, with migration name add_password_digest_to_users, appears as follows:

```
$ rails generate migration add_password_digest_to_users password_digest:string
```

Here we've also supplied the argument password_digest:string with the name and type of attribute we want to create. (Compare this to the original generation of the users table in [Listing 6.1](#), which included the arguments name:string and email:string.) By including password_digest:string, we've given Rails enough information to construct the entire migration for us, as seen in [Listing 6.35](#).

Listing 6.35: The migration to add a password_digest column.
db/migrate/[timestamp]_add_password_digest_to_users.rb

```
class AddPasswordDigestToUsers < ActiveRecord::Migration[5.0]
  def change
    add_column :users, :password_digest, :string
  end
end
```

[Listing 6.35](#) uses the add_column method to add a password_digest column to the users table. To apply it, we just migrate the database:

```
$ rails db:migrate
```

To make the password digest, has_secure_password uses a state-of-the-art hash function called [bcrypt](#). By hashing the password with bcrypt, we ensure that an attacker won't be able to log in to the site even if they manage to obtain a copy of the database. To use bcrypt in the sample application, we need to add the bcrypt gem to our Gemfile ([Listing 6.36](#)).²⁰

Listing 6.36: Adding bcrypt to the Gemfile.

```
source 'https://rubygems.org'

gem 'rails', '5.1.4'
gem 'bcrypt', '3.1.11'

.
```


Then run `bundle install` as usual:

```
$ bundle install
```

6.3.2 User has secure password

Now that we've supplied the User model with the required `password_digest` attribute and installed `bcrypt`, we're ready to add `has_secure_password` to the User model, as shown in [Listing 6.37](#).

Listing 6.37: Adding `has_secure_password` to the User model. **red** `app/models/user.rb`

```
class User < ApplicationRecord
  before_save { self.email = email.downcase }
  validates :name, presence: true, length: { maximum: 50 }
  VALID_EMAIL_REGEX = /\A[\w+\-\.]+\@[a-z\d\-\.]+\.[a-z]+\z/i
  validates :email, presence: true, length: { maximum: 255 },
    format: { with: VALID_EMAIL_REGEX },
    uniqueness: { case_sensitive: false }
  has_secure_password
end
```

As indicated by the **red** indicator in [Listing 6.37](#), the tests are now failing, as you can confirm at the command line:

Listing 6.38: **red**

```
$ rails test
```

The reason is that, as noted in [Section 6.3.1](#), `has_secure_password` enforces validations on the virtual password and `password_confirmation` attributes, but the tests in [Listing 6.26](#) create an `@user` variable without these attributes:

```
def setup
  @user = User.new(name: "Example User", email: "user@example.com")
end
```

So, to get the test suite passing again, we just need to add a password and its confirmation, as shown in [Listing 6.39](#).

Listing 6.39: Adding a password and its confirmation. **green** `test/models/user_test.rb`

```
require 'test_helper'

class UserTest < ActiveSupport::TestCase
  def setup
    @user = User.new(name: "Example User", email: "user@example.com",
      password: "foobar", password_confirmation: "foobar")
  end
  .
  .
  .
end
```

Note that the first line inside the `setup` method includes an additional comma at the end, as required by Ruby's hash syntax ([Section 4.3.3](#)). Leaving this comma off will produce a syntax error, and you should use your technical sophistication ([Box 1.1](#)) to identify and resolve such errors if (or, more realistically, when) they occur.

At this point the tests should be **green**:

Listing 6.40: **green**

```
$ rails test
```

We'll see in just a moment the benefits of adding `has_secure_password` to the User model ([Section 6.3.4](#)), but first we'll add a minimal requirement on password security.

Exercises

Solutions to exercises are available for free at railstutorial.org/solutions with any Rails Tutorial purchase. To see other people's answers and to record your own, join the [Learn Enough Society](http://learnenough.com/society) at learnenough.com/society.

1. Confirm that a user with valid name and email still isn't valid overall.
2. What are the error messages for a user with no password?

6.3.3 Minimum password standards

It's good practice in general to enforce some minimum standards on passwords to make them harder to guess. There are many options for [enforcing password strength in Rails](#), but for simplicity we'll just enforce a minimum length and the requirement that the password not be blank. Picking a length of 6 as a reasonable minimum leads to the validation test shown in [Listing 6.41](#).

Listing 6.41: Testing for a minimum password length. **red** test/models/user_test.rb

```
require 'test_helper'

class UserTest < ActiveSupport::TestCase
  def setup
    @user = User.new(name: "Example User", email: "user@example.com",
                     password: "foobar", password_confirmation: "foobar")
  end
  .
  .
  .
  test "password should be present (nonblank)" do
    @user.password = @user.password_confirmation = " " * 6
    assert_not @user.valid?
  end

  test "password should have a minimum length" do
    @user.password = @user.password_confirmation = "a" * 5
    assert_not @user.valid?
  end
end
```

Note the use of the compact multiple assignment

```
@user.password = @user.password_confirmation = "a" * 5
```

in [Listing 6.41](#). This arranges to assign a particular value to the password and its confirmation at the same time (in this case, a string of length 5, constructed using string multiplication as in [Listing 6.14](#)).

You may be able to guess the code for enforcing a minimum length constraint by referring to the corresponding maximum validation for the user's name ([Listing 6.16](#)):

```
validates :password, length: { minimum: 6 }
```

Combining this with a presence validation ([Section 6.2.2](#)) to ensure nonblank passwords, this leads to the User model shown in [Listing 6.42](#). (It turns out the `has_secure_password` method includes a presence validation, but unfortunately it only applies to records with *empty* passwords, which allows users to create invalid passwords like ' (six spaces).)

Listing 6.42: The complete implementation for secure passwords. **green** app/models/user.rb

```
class User < ApplicationRecord
  before_save { self.email = email.downcase }
  validates :name, presence: true, length: { maximum: 50 }
  VALID_EMAIL_REGEX = /\A[\w+\-\.]+\@[a-z\d\-\.\.]+\.[a-z]+\z/i
  validates :email, presence: true, length: { maximum: 255 },
    format: { with: VALID_EMAIL_REGEX },
    uniqueness: { case_sensitive: false }
  has_secure_password
  validates :password, presence: true, length: { minimum: 6 }
end
```

At this point, the tests should be **green**:

Listing 6.43: **green**

```
$ rails test:models
```

Exercises

Solutions to exercises are available for free at railstutorial.org/solutions with any Rails Tutorial purchase. To see other people's answers and to record your own, join the [Learn Enough Society](http://learnenough.com/society) at learnenough.com/society.

1. Confirm that a user with valid name and email but a too-short password isn't valid.
2. What are the associated error messages?

6.3.4 Creating and authenticating a user

Now that the basic User model is complete, we'll create a user in the database as preparation for making a page to show the user's information in [Section 7.1](#). We'll also take a more concrete look at the effects of adding `has_secure_password` to the User model, including an examination of the important `authenticate` method.

Since users can't yet sign up for the sample application through the web—that's the goal of [Chapter 7](#)—we'll use the Rails console to create a new user by hand. For convenience, we'll use the `create` method discussed in [Section 6.1.3](#), but in the present case we'll take care *not* to start in a sandbox so that the resulting user will be saved to the database. This means starting an ordinary rails console session and then creating a user with a valid name and email address together with a valid password and matching confirmation:

```
$ rails console
>> User.create(name: "Michael Hartl", email: "mhartl@example.com",
?>           password: "foobar", password_confirmation: "foobar")
=> #<User id: 1, name: "Michael Hartl", email: "mhartl@example.com",
created_at: "2016-05-23 20:36:46", updated_at: "2016-05-23 20:36:46",
password_digest: "$2a$10$xxucoRlMp06RLJSfWpZ8h08Dt9AZXlGRi3usP3njQg3...">
```

To check that this worked, let's look at the resulting users table in the development database using DB Browser for SQLite, as shown in [Figure 6.9](#).²¹ (If you're using the cloud IDE, you should download the database file as in [Figure 6.5](#).) Note that the columns correspond to the attributes of the data model defined in [Figure 6.8](#).

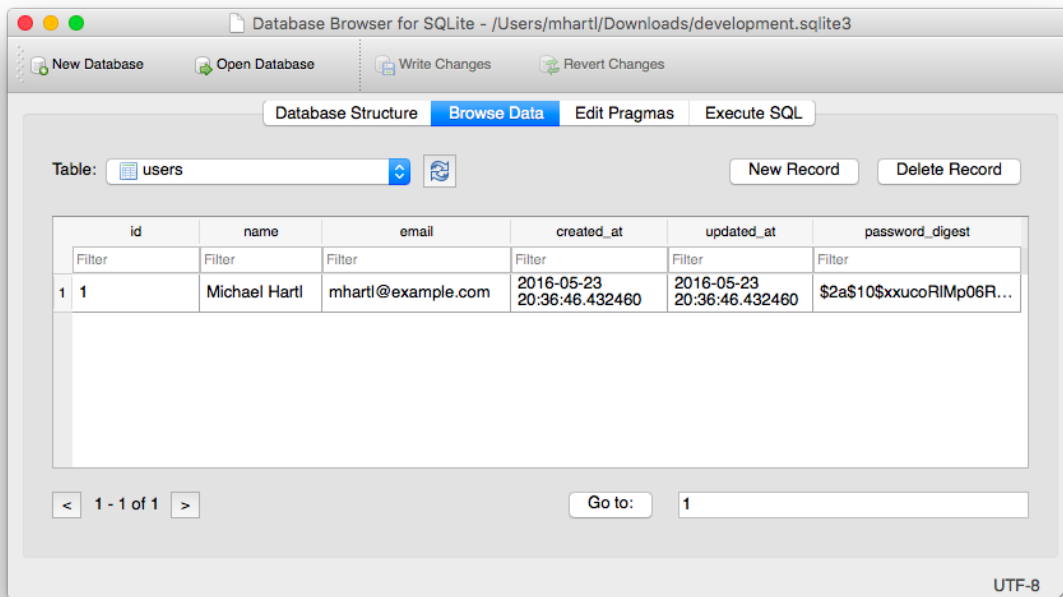


Figure 6.9: A user row in the SQLite database `db/development.sqlite3`.

Returning to the console, we can see the effect of `has_secure_password` from [Listing 6.42](#) by looking at the `password_digest` attribute:

```
>> user = User.find_by(email: "mhartl@example.com")
>> user.password_digest
=> "$2a$10$xxucoRlMp06RLJSfWpZ8h08Dt9AZXlGRi3usP3njQg3y0cVFzb6oK"
```

This is the hashed version of the password ("foobar") used to initialize the user object. Because it's constructed using `bcrypt`, it is computationally impractical to use the digest to discover the original password.²²

As noted in [Section 6.3.1](#), `has_secure_password` automatically adds an `authenticate` method to the corresponding model objects. This method determines if a given password is valid for a particular user by computing its digest and comparing the result to `password_digest` in the database. In the case of the user we just created, we can try a couple of invalid passwords as follows:

```
>> user.authenticate("not_the_right_password")
      false
>> user.authenticate("foobaz")
      false
```

Here `user.authenticate` returns `false` for invalid password. If we instead authenticate with the correct password, `authenticate` returns the user itself:

```
>> user.authenticate("foobar")
=> #<User id: 1, name: "Michael Hartl", email: "mhartl@example.com",
    created_at: "2016-05-23 20:36:46", updated_at: "2016-05-23 20:36:46",
    password_digest: "$2a$10$xxucoRlMp06RLJSfWpZ8h08Dt9AZXlGRi3usP3njQg3...">
```

In [Chapter 8](#), we'll use the `authenticate` method to sign registered users into our site. In fact, it will turn out not to be important to us that `authenticate` returns the user itself; all that will matter is that it returns a value that is `true` in a boolean context. Recalling from [Section 4.2.3](#) that `!!` converts an object to its corresponding boolean value, we can see that `user.authenticate` does the job nicely:

```
>> !!user.authenticate("foobar")
    => true
```

Exercises

Solutions to exercises are available for free at railstutorial.org/solutions with any Rails Tutorial purchase. To see other people's answers and to record your own, join the [Learn Enough Society](http://learnenough.com/society) at learnenough.com/society.

1. Quit and restart the console, and then find the user created in this section.
2. Try changing the name by assigning a new name and calling `save`. Why didn't it work?
3. Update user's name to use your name. *Hint*: The necessary technique is covered in [Section 6.1.5](#).

6.4 Conclusion

Starting from scratch, in this chapter we created a working `User` model with name, email, and password attributes, together with validations enforcing several important constraints on their values. In addition, we have the ability to securely authenticate users using a given password. This is a remarkable amount of functionality for only twelve lines of code.

In the next chapter, [Chapter 7](#), we'll make a working signup form to create new users, together with a page to display each user's information. In [Chapter 8](#), we'll then use the authentication machinery from [Section 6.3](#) to let users log into the site.

If you're using Git, now would be a good time to commit if you haven't done so in a while:

```
$ rails test
$ git add -A
$ git commit -m "Make a basic User model (including secure passwords)"
```

Then merge back into the master branch and push to the remote repository:

```
$ git checkout master
$ git merge modeling-users
$ git push
```

To get the `User` model working in production, we need to run the migrations at Heroku, which we can do with `heroku run`:

```
$ rails test
$ git push heroku
$ heroku run rails db:migrate
```

We can verify that this worked by running a console in production:

```
$ heroku run rails console --sandbox
>> User.create(name: "Michael Hartl", email: "michael@example.com",
```

```
?> password: "foobar", password_confirmation: "foobar")
=> #<User id: 1, name: "Michael Hartl", email: "michael@example.com",
created_at: "2016-05-23 20:54:41", updated_at: "2016-05-23 20:54:41",
password_digest: "$2a$10$74xFguZR0tZBXTUqs1Fjp0f30oLhrvgxC2wlohtTECH...">
```

6.4.1 What we learned in this chapter

- Migrations allow us to modify our application's data model.
 - Active Record comes with a large number of methods for creating and manipulating data models.
 - Active Record validations allow us to place constraints on the data in our models.
 - Common validations include presence, length, and format.
 - Regular expressions are cryptic but powerful.
 - Defining a database index improves lookup efficiency while allowing enforcement of uniqueness at the database level.
 - We can add a secure password to a model using the built-in `has_secure_password` method.
1. The name comes from the “[active record pattern](#)”, identified and named in *Patterns of Enterprise Application Architecture* by Martin Fowler. ↑
 2. Officially pronounced “ess-cue-ell”, though the alternate pronunciation “sequel” is also common. You can differentiate an individual author's preference by the choice of indefinite article: those who write “a SQL database” prefer “sequel”, whereas those who write “an SQL database” prefer “ess-cue-ell”. As you'll soon see, I prefer the latter. ↑
 3. By using an email address as the username, we open the possibility of communicating with our users at a future date ([Chapter 11](#) and [Chapter 12](#)). ↑
 4. Don't worry about exactly how the `t` object manages to do this; the beauty of *abstraction layers* is that we don't have to know. We can just trust the `t` object to do its job. ↑
 5. Officially pronounced “ess-cue-ell-ite”, although the (mis)pronunciation “sequel-ite” is also common. ↑
 6. The only exception is in [Section 14.3.3](#). ↑
 7. The timestamps are recorded in [Coordinated Universal Time](#) (UTC), which for most practical purposes is the same as [Greenwich Mean Time](#). From the [NIST Time and Frequency FAQ](#): **Q:** Why is UTC used as the acronym for Coordinated Universal Time instead of CUT? **A:** In 1970 the Coordinated Universal Time system was devised by an international advisory group of technical experts within the International Telecommunication Union (ITU). The ITU felt it was best to designate a single abbreviation for use in all languages in order to minimize confusion. Since unanimous agreement could not be achieved on using either the English word order, CUT, or the French word order, TUC, the acronym UTC was chosen as a compromise. ↑
 8. Exceptions and exception handling are somewhat advanced Ruby subjects, and we won't need them much in this book. They are important, though, and I suggest learning about them using one of the Ruby books recommended in [Section 14.4.1](#). ↑
 9. The `update_attributes` method is an alias for the shorter `update` method, but I prefer the longer version because of its similarity to the singular version of the method, `update_attribute`. ↑
 10. I'll omit the output of console commands when they are not particularly instructive—for example, the results of `User.new`. ↑
 11. For example, did you know that “Michael Hartl@example.com, with quotation marks and a space in the middle, is a valid email address according to the standard? Incredibly, it is—but it's absurd. ↑
 12. Note that, in [Table 6.1](#), “letter” really means “lower-case letter”, but the `i` at the end of the regex enforces case-insensitive matching. ↑
 13. If you find it as useful as I do, I encourage you to [donate to Rubular](#) to reward developer [Michael Lovitt](#) for his wonderful work. ↑
 14. As noted briefly in the introduction to this section, there is a dedicated test database, `db/test.sqlite3`, for this purpose. ↑
 15. Technically, only the domain part of the email address is case-insensitive: `foo@bar.com` is actually different from `Foo@bar.com`. In practice, though, it is a bad idea to rely on this fact; as noted at [about.com](#), “Since the case sensitivity of email addresses can create a lot of confusion, interoperability problems and widespread headaches, it would be foolish to require email addresses to be typed with the correct case. Hardly any email service or ISP does enforce case sensitive email addresses, returning messages whose recipient's email address was not typed correctly (in all upper case, for example).” Thanks to reader Riley Moses for pointing this out. ↑
 16. Of course, we could just edit the migration file for the `users` table in [Listing 6.2](#), but that would require rolling back and then migrating back up. The Rails Way™ is to use migrations every time we discover that our data model needs to change. ↑
 17. See the [Rails API entry on callbacks](#) for more information on which callbacks Rails supports. ↑
 18. In this context, *virtual* means that the attributes exist on the model object but do not correspond to columns in the database. ↑
 19. Hashed password digests are often erroneously referred to as *encrypted passwords*. For example, the [source code](#) of `has_secure_password` makes this mistake, as did the first two editions of this tutorial. This terminology is wrong because by design encryption is *reversible*—the ability to encrypt implies the ability to

decrypt as well. In contrast, the whole point of calculating a password's hash digest is to be *irreversible*, so that it is computationally intractable to infer the original password from the digest. (Thanks to reader Andy Philips for pointing out this issue and for encouraging me to fix the broken terminology.) ↑

20. As always, you should use the version numbers listed at gemfiles-4th-ed.railstutorial.org instead of the ones listed here. ↑
21. If for any reason something went wrong, you can always reset the database as follows:
 1. Quit the console.
 2. Run `$ rm -f development.sqlite3` at the command line to remove the database. (We'll learn a more elegant method for doing this in [Chapter 7](#).)
 3. Re-run the migrations using `$ rails db:migrate`.
 4. Restart the console. ↑
22. By design, the bcrypt algorithm produces a [salted hash](#), which protects against two important classes of attacks ([dictionary attacks](#) and [rainbow table attacks](#)). ↑